

An Overview of R

Nathan Eastwood

04/03/2019

Let's Connect!



nathaneastwood



@nathaneastwood_



nathaneastwood



<http://nathaneastwood.github.io> (<http://nathaneastwood.github.io>)

About Me

Education

University	Degree	Grade
Plymouth University	BSc Mathematics and Statistics	1 st Class (Hons)
Sheffield University	MSc Statistics	Distinction

Employment History

Company	Title	Time
Nuffield Health Foundation	Data Analyst	Summer 2011
C3 Resources	Data Analyst	Jun. 2012 - Jan. 2013
Plymouth University	Data Scientist	Feb. 2013 - Feb. 2016
Mango Solutions	Data Science Consultant	Mar. 2016 - May 2018
iotec Global	Senior Data Scientist	May 2018 - Aug. 2018
Equiniti Data	Senior Data Scientist	Aug. 2018 - Present

Clients Worked With

- Ministry of Defence
- Public Health England
- Office for National Statistics
- HAYS Recruitment
- NATO
- Direct Line Group
- and many more!

Key Skills

Language

Years of Experience

R

9 years

Python

4 years

SQL

7 years

LaTeX

7 years

bash

6 years

git

7 years

My Type of Work

- Software development
- Building data analysis pipelines
- Data analysis
- Machine learning
- Statistical modelling
- Data visualisation

A Brief History of R

The Birth of R

- Before R, there was S
- S was created at Bell Labs in 1976 by John Chambers
- S Plus was developed by TIBCO Software in 1988
- R was created in 1992 by Ross Ihaka and Robert Gentleman at the University of Auckland, New Zealand
- The first stable release was in 1995
- R is now maintained by the R Development Core Team

CRAN

- The Comprehensive R Archive Network (CRAN) is where you will find the majority (over 13,500) supplementary packages
- You can find additional packages on Bioconductor, Omegahat, GitHub, etc.

R and Other Languages

- R is written with a combination of C and Fortran
- More advanced users can write C, C++, .NET, Java, Python, JavaScript and Go code and call it directly from R

The Birth of Data Science

- R has become increasingly popular. As of February 2019, R ranks 15th in the TIOBE index¹, a measure of popularity of programming languages
- 90% of the world's data was created in the last two years²

¹: <https://www.tiobe.com/tiobe-index/>

²: <https://www.mediapost.com/publications/article/291358/90-of-todays-data-created-in-two-years.html>

RStudio

- RStudio are a company which develop the fantastic (and free) R IDE of the same name
- The initial release was in 2011
- RStudio and its team have contributed to many open source packages including
 - Tidyverse – R packages for data science, including `ggplot2`, `dplyr`, `tidyr`, and `purrr`
 - `shiny` – An interactive web technology
 - RMarkdown – Insert R code into markdown documents
 - `knitr` – Dynamic reports combining R, TeX, Markdown & HTML
 - `packrat` – Package dependency tool
 - `devtools` – Package development tool

What Can R Do?

Interpreted Language

```
2 + 2
```

```
[1] 4
```


Data Manipulation

readr

```
qof_data <- read_csv(  
  here(  
    'data',  
    'QOF 2017-18- Prevalence, achievements and exceptions at CCG level.csv'  
  )  
)
```

readr

qof_data

dplyr

dplyr

- dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges
 - `mutate()` adds new variables that are functions of existing variables
 - `select()` picks variables based on their names.
 - `filter()` picks cases based on their values.
 - `summarise()` reduces multiple values down to a single summary.
 - `arrange()` changes the ordering of the rows.
- These all combine naturally with `group_by()` which allows you to perform any operation “by group”

Renaming Columns

```
qof_data <-  
  qof_data %>%  
    rename_all(funs(gsub("\\(|\\)|\\+", "", gsub(" |-", "_", tolower(.)))))  
qof_data
```

Selecting Columns

```
alt_data <-
  qof_data %>%
  select(
    ccg_name, total_exceptions_2016_2017, total_exceptions_2017_2018,
    patients_receiving_intervention_per_cent
  )
alt_data
```

```
# A tibble: 194 x 4
```

	ccg_name <chr>	total_exceptions_2016_2017 <dbl>	total_exceptions_2017_2018 <dbl>	patients_receiving_intervention_per_cent <dbl>
1	NHS AIREDALE, W...	328	342	89.5
2	NHS BARNSELEY CCG	227	241	91.9
3	NHS BASSETLAW C...	71	58	90.9
4	NHS BRADFORD DI...	519	480	89.7
5	NHS CALDERDALE ...	181	177	90.8
6	NHS BRADFORD CI...	110	264	93.8
7	NHS DONCASTER C...	482	572	91.1
8	NHS EAST RIDING...	481	539	90.4
9	NHS GREATER HUD...	494	492	91.5
10	NHS HAMBLETON, ...	118	165	91.3

```
# ... with 184 more rows
```

Filtering Columns

```
alt_data <-
  alt_data %>%
  filter(ccg_name %in% c(
    "NHS KERNOW CCG", "NHS NORTHERN, EASTERN AND WESTERN DEVON CCG",
    "NHS SOUTH DEVON AND TORBAY CCG"
  ))
alt_data
```

```
# A tibble: 3 x 4
  ccg_name                total_exceptions_2... total_exceptions_... patients_receiving_i...
  <chr>                  <dbl>                <dbl>                <dbl>
1 NHS KERNOW CCG          1941                2183                90.8
2 NHS NORTHERN, EA...     2050                2885                90.7
3 NHS SOUTH DEVON ...     785                 855                89.7
```


Sorting Columns

```
alt_data <-  
  alt_data %>%  
    arrange(patients_receiving_intervention_per_cent)  
alt_data
```

```
# A tibble: 3 x 4  
  ccg_name                total_exceptions_2... total_exceptions_... patients_receiving_i...  
  <chr>                  <dbl>                <dbl>                <dbl>  
1 NHS SOUTH DEVON ...      785                  855                  89.7  
2 NHS NORTHERN, EA...    2050                 2885                 90.7  
3 NHS KERNOW CCG        1941                 2183                 90.8
```

Mutating Columns

```
alt_data <-
  alt_data %>%
  mutate(
    exception_rate_yoy_change = (total_exceptions_2016_2017 / total_exceptions_2017_2018) * 100
  )
alt_data
```

```
# A tibble: 3 x 5
  ccg_name total_exceptions... total_exception... patients_receiv... exception_rate_...
  <chr>          <dbl>          <dbl>          <dbl>          <dbl>
1 NHS SOUT...      785            855            89.7            91.8
2 NHS NORT...    2050           2885            90.7            71.1
3 NHS KERN...    1941           2183            90.8            88.9
```

dplyr Chaining

```
gof_data %>%
  filter(ccg_name %in% c(
    "NHS KERNOW CCG", "NHS NORTHERN, EASTERN AND WESTERN DEVON CCG",
    "NHS SOUTH DEVON AND TORBAY CCG"
  )) %>%
  select(
    ccg_name, total_exceptions_2016_2017, total_exceptions_2017_2018,
    patients_receiving_intervention_per_cent
  ) %>%
  arrange(patients_receiving_intervention_per_cent) %>%
  mutate(
    exception_rate_yoy_change = (total_exceptions_2016_2017 / total_exceptions_2017_2018) * 100
  )
```

A tibble: 3 x 5

	ccg_name	total_exceptions...	total_exception...	patients_receiv...	exception_rate_...
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	NHS SOUT...	785	855	89.7	91.8
2	NHS NORT...	2050	2885	90.7	71.1
3	NHS KERN...	1941	2183	90.8	88.9

Summarising Data

```
gof_data %>%  
  summarise(  
    mean_prip = mean(patients_receiving_intervention_per_cent),  
    max_prip = max(patients_receiving_intervention_per_cent),  
    min_prip = min(patients_receiving_intervention_per_cent)  
  ) %>%  
  mutate(range_prip = max_prip - min_prip)
```

```
# A tibble: 1 x 4  
  mean_prip max_prip min_prip range_prip  
    <dbl>    <dbl>    <dbl>    <dbl>  
1     90.6     94.0     82.6     11.3
```

Grouping Data

```
gof_data %>%
  group_by(sub_region_name) %>%
  summarise(mean_prip = mean(patients_receiving_intervention_per_cent))
```

```
# A tibble: 16 x 2
```

sub_region_name	mean_prip
<chr>	<dbl>
1 LONDON NORTH EAST AND CENTRAL	91.6
2 LONDON NORTH WEST	88.0
3 LONDON SOUTH	88.7
4 NHS ENGLAND CENTRAL MIDLANDS	91.2
5 NHS ENGLAND CHESHIRE AND MERSEYSIDE	91.1
6 NHS ENGLAND CUMBRIA AND NORTH EAST	90.8
7 NHS ENGLAND EAST	90.9
8 NHS ENGLAND GREATER MANCHESTER	90.9
9 NHS ENGLAND HAMPSHIRE, ISLE OF WIGHT AND THAMES VALLEY	89.7
10 NHS ENGLAND KENT, SURREY AND SUSSEX	90.3

```
# ... with 6 more rows
```

Plotting

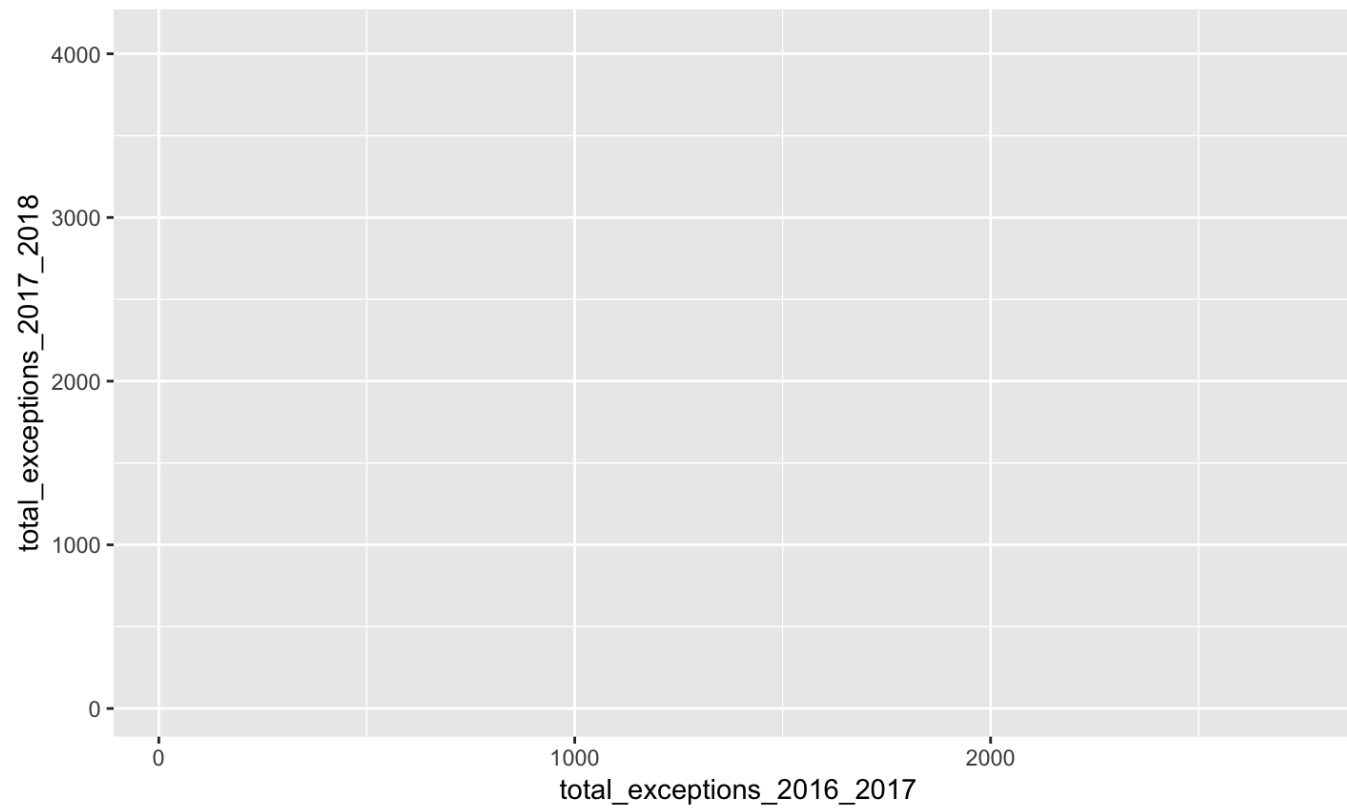
History of Plotting in R

- R comes with the base `graphics` package
- Later the `lattice` package and `grid` package were developed
- `ggplot2` was developed on top of `grid` and is developed on the theory of the “grammar of graphics”

ggplot2

```
p <- ggplot(  
  data = gof_data,  
  aes(x = total_exceptions_2016_2017, y = total_exceptions_2017_2018)  
)
```

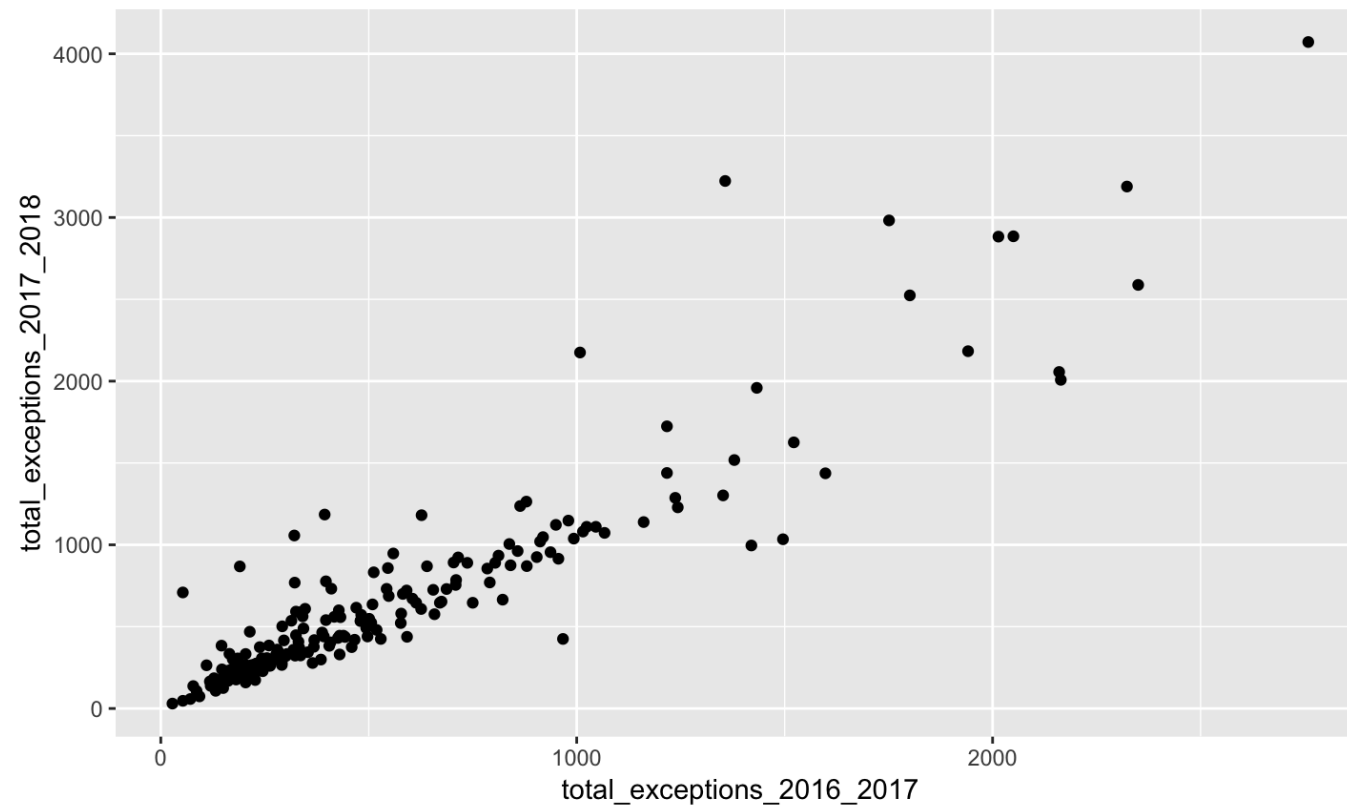

ggplot2



Points

```
p <- p +  
  geom_point()
```


Points



Lines

```
p <- p +  
  geom_smooth(method = "lm")
```

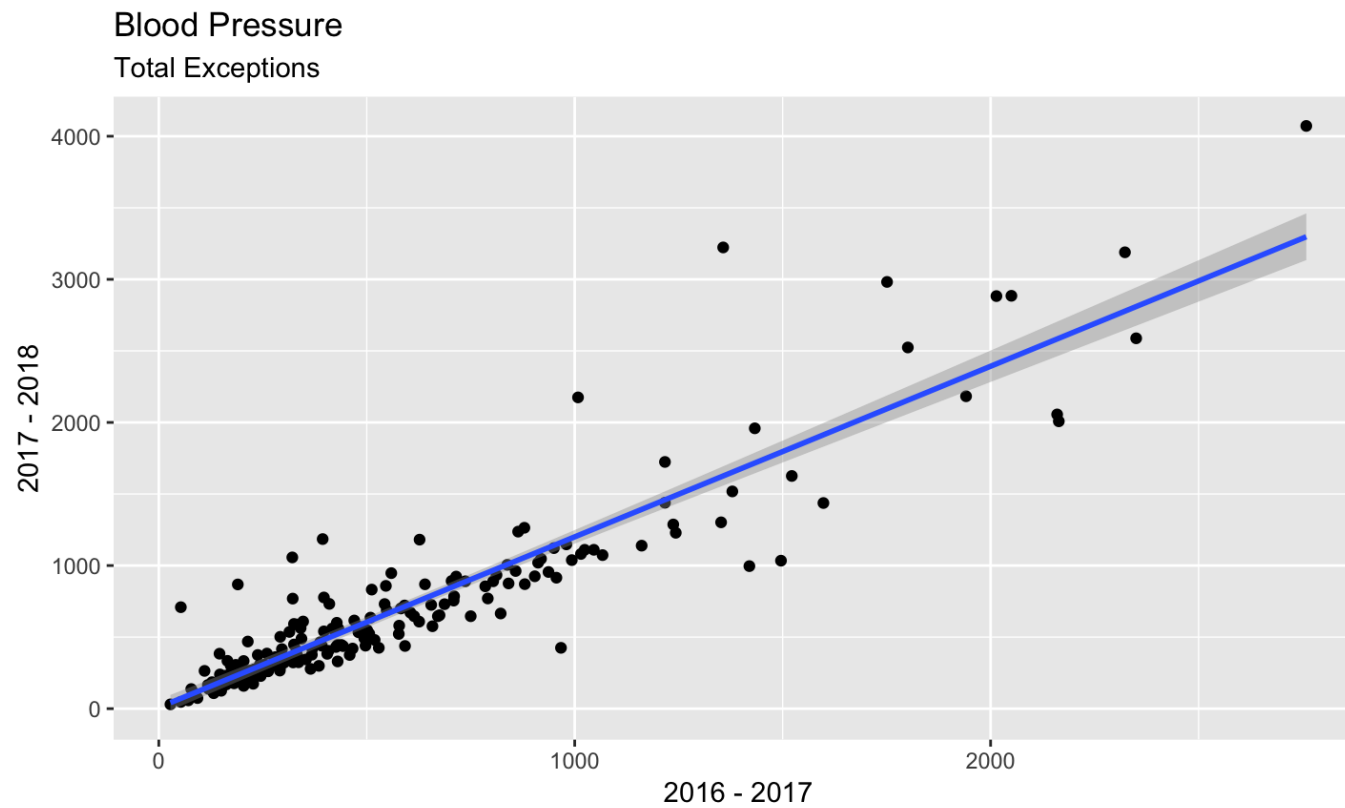
Lines



Labels

```
p <-  
  p +  
  labs(  
    x = "2016 - 2017", y = "2017 - 2018",  
    title = "Blood Pressure", subtitle = "Total Exceptions"  
  )
```

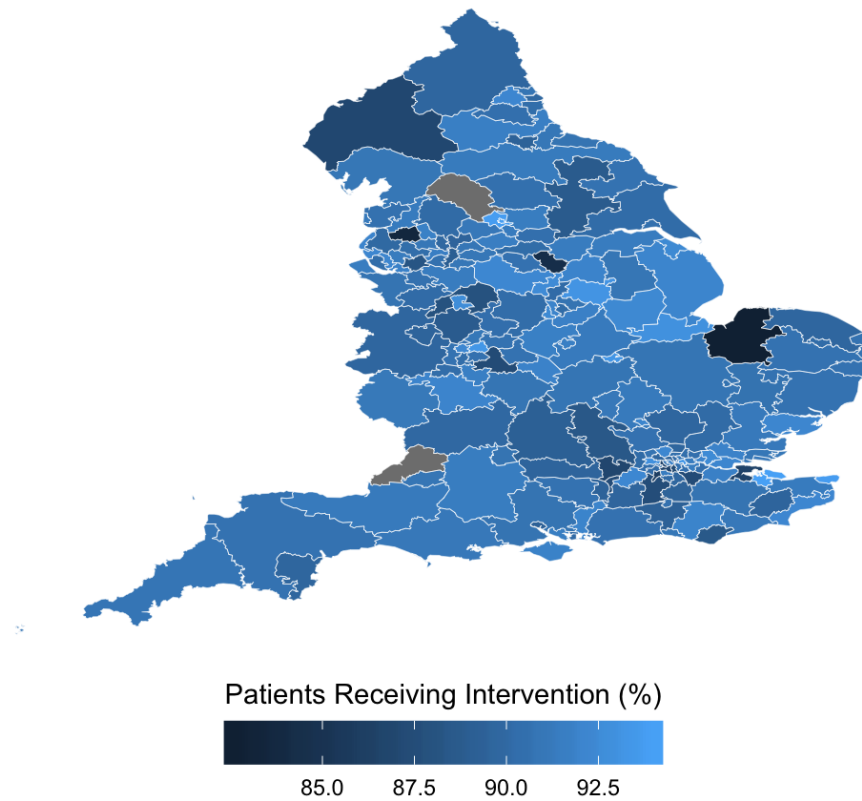
Labels




ggplot2 Extensions

- There are many extension packages that have been written for `ggplot2`
- These packages are showcased at <https://www.ggplot2-exts.org/>
(<https://www.ggplot2-exts.org/>)

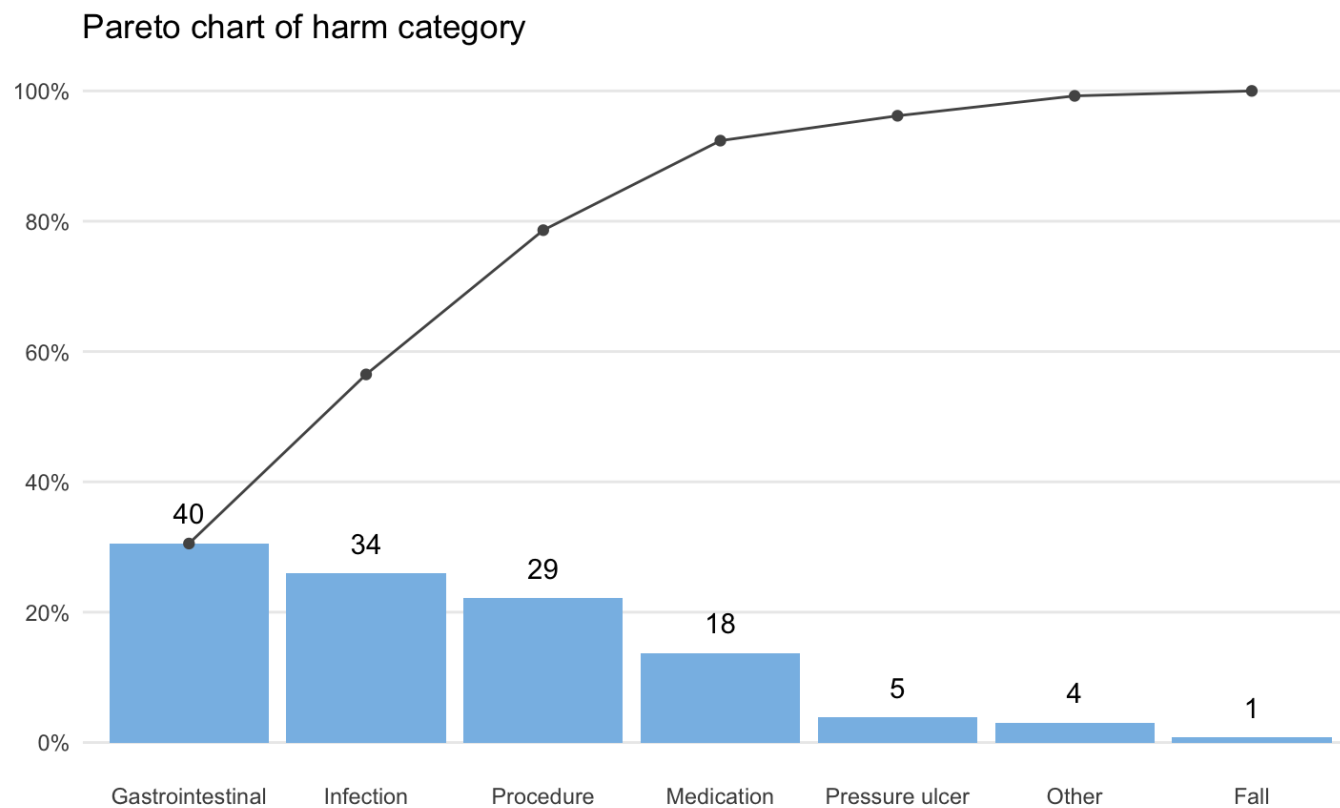
ggplot2 Map



qicharts2



Pareto charts



Statistical Programming

Statistical Programming

- R is first and foremost designed as a statistical programming language
- It can do all sorts of statistical analyses, from t -tests to machine learning
- There are hundreds of advanced, production quality packages for statistical analysis such as `caret` and `mlr`

Linear Regression Example

```
lm1 <- lm(total_exceptions_2016_2017 ~ total_exceptions_2017_2018, data = gof_data)
summary(lm1)
```

So R Is Just a Statistical Programming Language?

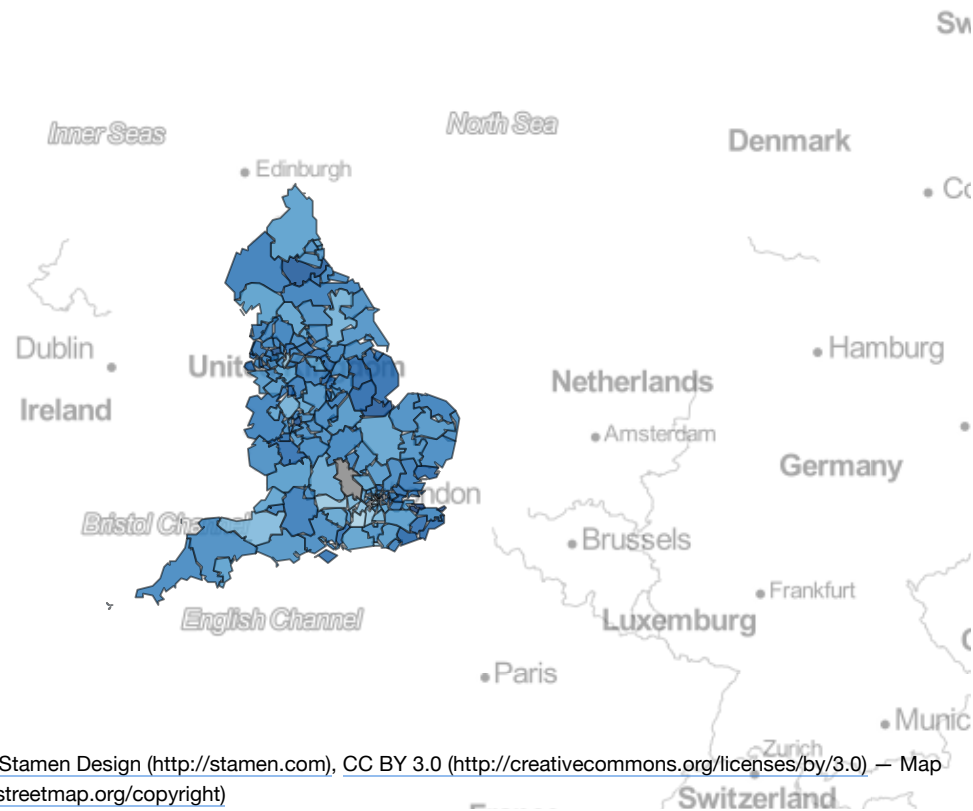
Not Quite!

- R can...
 - Make interactive graphics
 - Make web applications
 - Produce professional, word, pdf, excel, etc. documents
 - Communicate with other languages such as Java, C++, Python and Go!

Interactive Plotting

- You can create interactive graphics using `htmlwidgets`
- `htmlwidgets` provides a wrapper for JavaScript libraries
- You can see `htmlwidgets` examples at http://www.htmlwidgets.org/showcase_leaflet.html
(http://www.htmlwidgets.org/showcase_leaflet.html)

Leaflet example



Leaflet (<http://leafletjs.com>) | Map tiles by Stamen Design (<http://stamen.com>), CC BY 3.0 (<http://creativecommons.org/licenses/by/3.0>) — Map data © OpenStreetMap (<http://www.openstreetmap.org/copyright>)

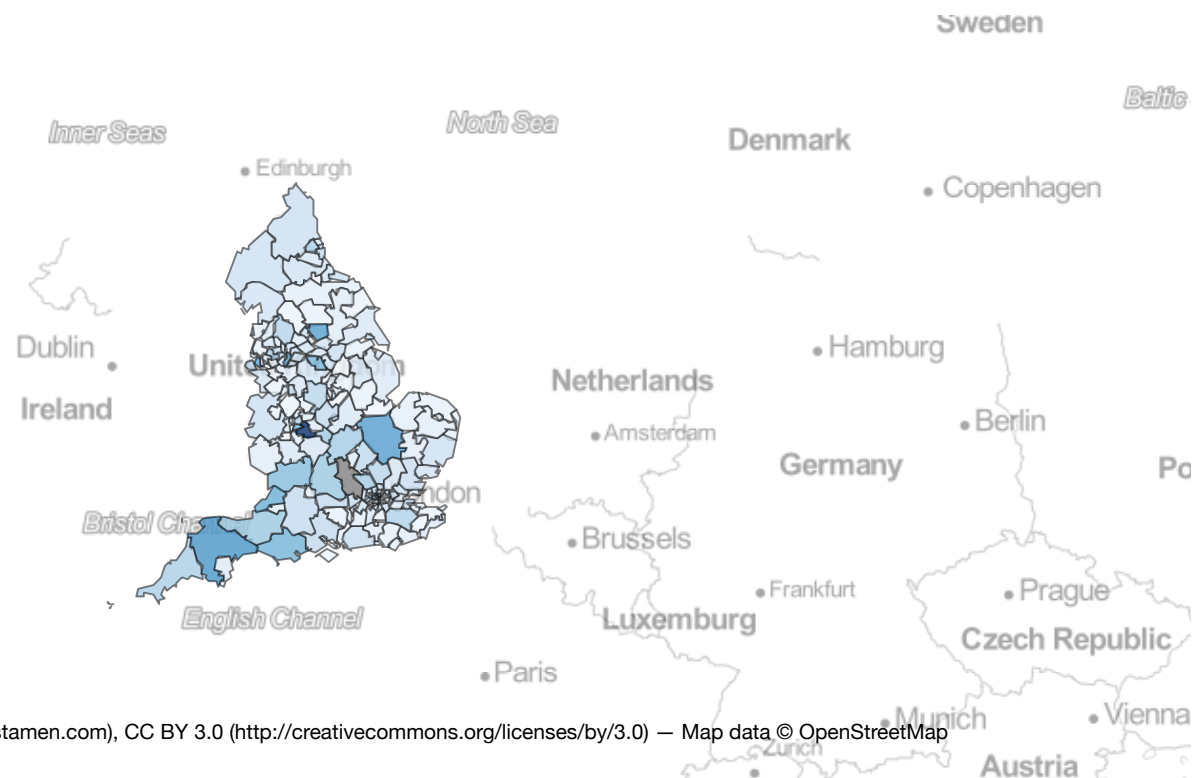
Application Building

- It is possible to build interactive web applications in R by using `shiny`
- `shiny` wraps HTML, JavaScript and CSS into an R package which you call using only R code
- You can use custom HTML, JavaScript and CSS to extend `shiny`
- There are lots of `shiny` application examples at <https://shiny.rstudio.com/gallery/> (<https://shiny.rstudio.com/gallery/>)

Shiny Example

Statistic:

Number of practices 2016_2017



Leaflet (<http://leafletjs.com>) | Map tiles by Stamen Design (<http://stamen.com>), CC BY 3.0 (<http://creativecommons.org/licenses/by/3.0>) — Map data © OpenStreetMap (<http://www.openstreetmap.org/copyright>)

Documents

- You can build many documents from R using `Rmarkdown`
- This is an `Rmarkdown` presentation!
- Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents (and more!)
- When you **knit** together an `Rmarkdown` document, a document will be generated that includes both content as well as the output of any embedded R code “chunks” within the document
- For more details on using R Markdown see <http://rmarkdown.rstudio.com>
(<http://rmarkdown.rstudio.com>)

R can...

- Read in data
- Clean and manipulate data
- Produce complex graphics
- Perform statistical analyses
- Create interactive graphics
- Create websites and web applications
- Produce documents of many formats

Any Questions?



nathaneastwood



@nathaneastwood_



nathaneastwood



<http://nathaneastwood.github.io> (<http://nathaneastwood.github.io>)